
Electrical Engineering and Computer Science
EECS 358 - INTRODUCTION TO PARALLEL COMPUTING

Lecture 1
Parallel Processing

Outline

- What is parallel processing
- Why parallel processing
- Parallel processing options
- Aspects of parallel processing
 - Architectures
 - Software
 - Algorithms
- Summary
- READING : Kumar - ch 1,2
- Web: www.intel.com, www.ibm.com

Course Administration

- Electrical Engineering and Computer Science 358: Introduction to Parallel Computing
- Class meets Tuesdays and Thursdays 2:00-3:20PM
- Instructor:
 - Gokhan Memik
 - Office: L475 Tech, Phone: (847) 467-1168,
 - Email: memik@eecs.northwestern.edu,
 - Web: <http://www.eecs.northwestern.edu/~memik>
 - Office Hours: Thursdays, 3:30-4:30pm (or by appointment)
- Teaching Assistant:
 - Majed Beigi
 - Room: Tech L458
 - Email: majed.beigi@northwestern.edu

- Office Hours: TBD
- Grader:
 - Wenting Zhou
 - Room: Tech L458
 - Email: wentingzhou2017@u.northwestern.edu
- Class Material:
 - Everything on Canvas one hour before the class

Outline of course

- Introduction to Parallel Programming (2 lectures)
- Shared Memory Parallel Programming (4 lectures)
- Distributed memory message-passing parallel programming (5 lectures)
- Data parallel and SIMD parallel programming (3 lectures)
- Parallel programming tools (1 lectures)
- Parallel algorithms and applications (3 lectures)

Books and Reading Material

Textbook:

- A. Grama, A. Gupta, G. Karypis, and Vipin Kumar, *Introduction to Parallel Computing*, Addison Wesley, 2nd edition, 2003

Other Material:

- S. Brawer, "Introduction to Parallel Programming", Academic Press, 2014.
- I. Foster, "Designing and Building Parallel Programs," Addison Wesley, 1994. Available online at: <http://www.mcs.anl.gov/dbpp>.
- B. Bauer, "Practical Parallel Programming," Academic Press, 1992.
- W. Gropp et al, "Using MPI: Portable Parallel Programming with the Message Passing Interface," MIT Press, 1994. Available online at: <http://www.mcs.anl.gov/mpi/index.html>.
- I will be using Canvas for course material

Grades for Class

- Three or four homeworks and programming assignments - 50 % total
- Two exams (25 % each) - 50 % total
Midterm (Late April/Early May)
Endterm (Last week of classes)
- There is a late penalty of 10% per calendar day and no homeworks/assignments will be accepted once the solutions are posted to the class web site.

What is Parallel Processing ?

- Parallel processing is the process by which a problem is solved using multiple resources working concurrently and collaboratively.
- Daily life examples:
 - House construction
 - Car manufacturing
 - Grocery store operation
- Effective parallel processing requires:
 - Appropriate division of labor
 - Proper coordination

Why Parallel Processing ?

- Parallel Processing is an effective answer for the tremendous future computing requirements:
 - Very complex applications
 - Real-time requirements of applications
- Existing algorithms for many applications running on uniprocessors inadequate for the future:
 - Need more variables for *better quality* results
 - Need to solve *larger problem sizes*
 - For many applications *fast-response is required*
- Large set of computational problems are inherently parallel in nature

Why Parallel Processing ?

- Computational Science Grand Challenges:
 - Science today: Experimentation, Theory, Simulation
 - Simulation relies heavily on parallel processing
- Aerodynamics: design of new aircraft
- Biology: modeling of genetic compounds
- Computer Science: VLSI CAD
- Physics: simulation of fusion reactions
- Civil Engineering: structural analysis
- Electrical Engineering: image and signal processing

Why Parallel Processing ?

- Rapidly growing parallel processing technology has recently become commercially available
 - No "uniprocessor" left
 - Most improvements in the future will come from parallel programming, in other words, everyone relies on parallel computing
- It is relatively easy to build hardware for systems with hundreds or thousands of processors which can theoretically achieve Teraflops/Petaflops
- Parallel processing has become more economical in recent years
 - 1 CRAY T90 processor - 2 GFlops (\$2,500,000)
 - 8 Node IBM SP2 using PowerPC processors - 2 GFlops (\$200,000)
 - 8 Node SGI Origin 2000 using R10000 processors - 10 GFlops (\$100,000)
 - Single chip IBM Cell with 8 cores - 20 to 200 GFlops (\$1,000)
 - Nvidia GeForce Titan - 1 to 10 TFlops (\$1,000)

Parallel Processing Options

- Chip multiprocessors (IBM Cell, Sun Niagara, Intel N-Core)
- Network of workstations (lowest cost)
- Multiprocessor workstations (\$15,000)
 - SUN Ultra 10, HP J5000, SGI
- Shared memory multiprocessors servers (\$10,000-\$200,000)
 - SGI Origin 2000, SUN Enterprise 6500, HP K580
- Distributed memory multicomputers (\$200,000-\$800,000)
 - IBM SP2
- Massively Parallel Processors (\$10,000,000)
 - SGI ASCI Blue, Tera Corporation MTA, IBM Blue Gene, NEC Earth Simulator, K Computer, IBM Sequoia, Cray Titan, Tianhe

Aspects of Parallel Processing

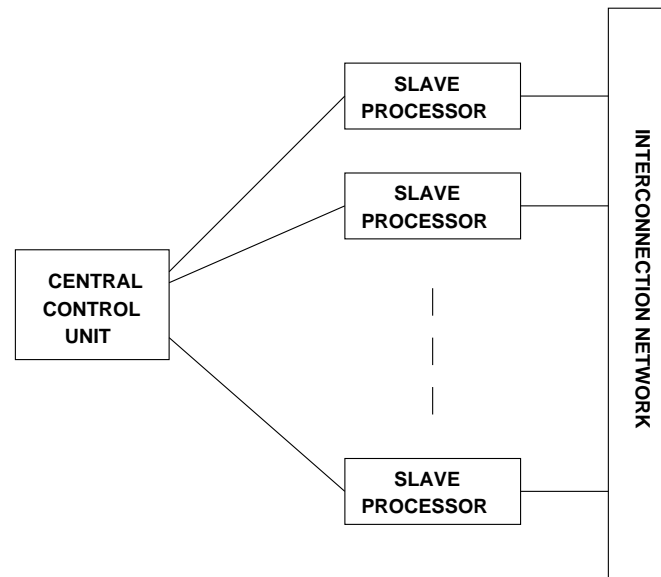
- Architectures:
 - Processors and memories connected together
- Software:
 - Operating systems
 - Compilers
 - Libraries
 - Tools - debuggers, performance analyzers
- Algorithms:
 - Designing software that best fits underlying architecture

Architecture

- Basic components of any architecture:
 - Processors and Memory (Processing Units)
 - Interconnect
- Logical classification based on:
 - Control mechanism - SIMD and MIMD
 - * SIMD (Single Instruction Multiple Data stream)
 - * MIMD (Multiple Instruction Multiple Data stream)
 - Address space organization - Shared and Distributed

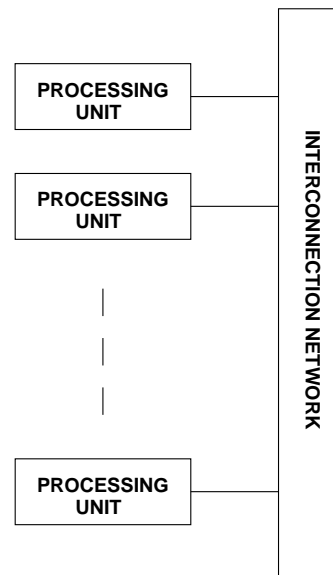
SIMD Architectures

- Centralized control processor defines the sequence of instructions executed by slaves
- Slave processors execute same instruction streams on distinct data sets



MIMD Architectures

- Each processor executes program independent of other processors




SIMD vs. MIMD

- SIMD:
 - Need custom hardware for processors
 - Slave processors are simple and therefore low cost
 - Good for programs with lot of synchronization due to lock step operation
- MIMD:
 - Easy to build out of commodity parts
 - Processors are complex, but availability of low cost commodity microprocessors offsets the SIMD advantage
 - Can handle a more general class of problems with reasonable efficiency


Shared Address Space

- Shared address space:
 - Processors can directly access all the data in the system
 - Interaction via modification of objects in shared space



Distributed Address Space

- Distributed Address space:
 - Processors can directly access only local data
 - Interaction via explicit message passing






Shared Address vs. Distributed Address

- Shared Address:
 - Transparent access to all program data makes them easier to program
 - On large shared address machines, performance issues similar to distributed address space machines
 - Debugging difficult due to "race" conditions
- Distributed Address:
 - Explicit fetching of non-local data makes them difficult to program
 - Model most suitable for constructing massively parallel machines


Static Interconnects

- Consist of point-to-point links between processors
- Can make parallel system expansion easy
- Some processors may be "closer" than others
- Examples : Hypercube, Mesh/Torus, Fat tree



Dynamic Interconnects

- Paths are established as needed between processors
- System expansion difficult
- Processors are usually equidistant
- Examples : Bus based, Crossbar, Multistage networks



Examples

- SGI Power Challenge
 - MIMD, Shared address space
 - Dynamic interconnect (bus)
- IBM SP-2
 - MIMD, Distributed address space
 - Dynamic interconnect (multistage)
- IBM Sequoia
 - MIMD, Distributed address space
 - Custom interconnect (Mesh/Torus)

Operating Systems

- Need to support tasks similar to serial OS like UNIX
 - Memory and process management, File systems, Security
- Additional support needed
 - Job scheduling - Time sharing, Space sharing
 - Parallel programming support - message passing, synchronization

Shared Address Space MIMD

- Typically time shared - many user programs co-exist in the system and compete for time slices
- Access to job queue can be centralized or decentralized
- Parallel programming support via primitives for creation, destruction and synchronization of processes

Distributed Address Space MIMD

- Usually access to parallel machine is via a host computer running a serial OS
- Nodes of parallel machine have a simpler version of OS
- Typically space shared - user programs have processor sub-sets allocated for exclusive use
- Parallel programming support via primitives for message passing and network management

Compilers

- Automatic parallelization:
 - Compiler detects and exploits concurrency optimally
 - Requires complex analysis and transformations
- Implicit parallel programming:
 - Programmer directives to help compiler analysis
 - Compiler performs all the "grunt" work
- Explicit parallel programming : easiest to support

Libraries

- Make using parallel machines easier
- Library implementations are usually done by skilled and experienced programmers working closely with machine designers resulting in high levels of performance.
- Library routines can be used as building blocks for complex applications
- Usually cover certain specialized application domains
- Examples : SCALAPACK, MPI, PVM

Tools

- Essential due to high degree of complexity in parallel machines
- Performance analyzers:
 - Help in identifying bottlenecks
 - Can identify relative importance of different parts of program with respect to possible performance gains
- Debugging:
 - Need to capture the state of multiple processes
 - Bugs commonly caused by synchronization errors are difficult to capture
- Examples : PARAGRAPH/PICL, Prism, Pablo, Upshot

Algorithms

- Divide a given problem into sub-problems
- Sub-problems must have a high degree of concurrency
- Must minimize synchronization among sub-problems
- Performance is a key aspect of algorithm design

Summary

- What is parallel processing
- Why parallel processing
- Parallel processing options
- Aspects of parallel processing
 - Architectures
 - Software
 - Algorithms
- NEXT CLASS : Introduction to Parallel Programming
- READING : Kumar - ch 1,2