

---

**Electrical Engineering and Computer Science**  
**EECS 358 - INTRODUCTION TO PARALLEL COMPUTING**

**Lecture 4**  
**Shared Memory Programming I**

---

# Outline


---

- An overview of shared memory parallel programming
- Process model: creation and destruction
- Shared variables
- Synchronization: locks, barriers
- READING: Kumar, Chapter 7

# Processes

---

- Processes are central to shared memory programming
- A process is a program along with all of **its enviroment** (support structures)
- A set of processes can share variables among themselves and therefore communicate
- All processes have a unique ID associated with them



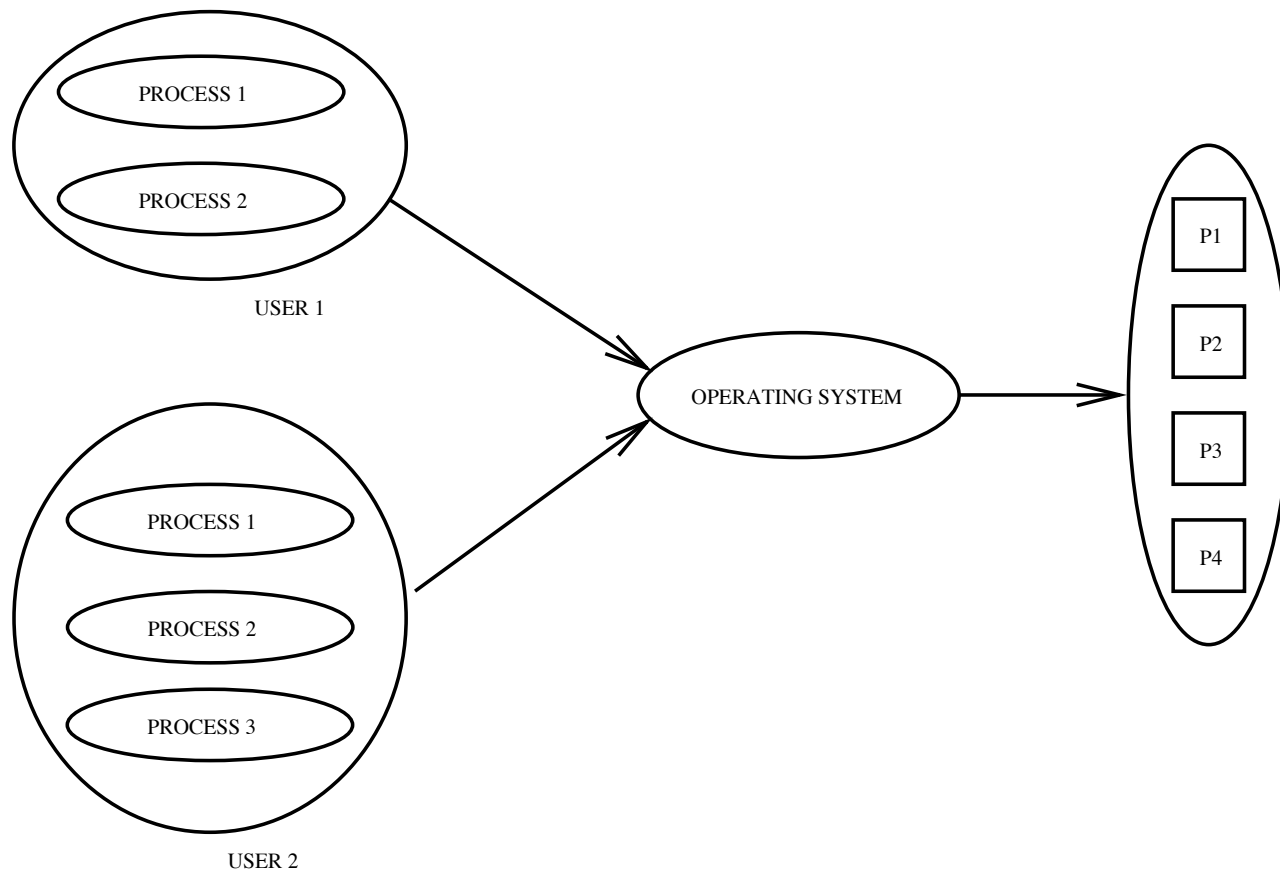
# Processes

---

- Processes are created at the user's request by the operating system
- Processes are managed entirely by the operating system
- Processes are mapped to the processors in the system by the operating system
- Mapping of processes to processors is not under user control - it is assumed to be random for the purposes of parallel programming

# Processes

---



# Processes

---


- Program execution is started using a single process (parent process)
- This process can create/destroy additional processes during program execution (child processes)
- Program computation is shared among parent and child processes:
  - Data parallel or function parallel
  - Coarse-grain or fine-grain
- Parallel execution achieved through mapping the processes to available physical processors in the system

# Process Creation

---

- We will illustrate parallel programming calls using SGI IRIX operating system, other machines similar
- Use a process creation primitive:

C calling conventions  
`m_set_procs(nproc);`  
`m_fork(func, [args]);`



# Process Creation

---

- Each child process is an exact copy of its parent
- All the variables associated with the parent are private for the parent and each child unless explicitly made shared
- The parent has ID=0, children share the other nproc-1 IDs

```
id = m_get_myid();  
nprocs = m_get_procs();
```




# Process Destruction

---

- Use a process destruction primitive

```
m_kill_procs();
```



# Process Destruction

---

- All child processes are terminated; only the parent remains
- The parent waits for all child processes to terminate before returning

# Example

---

```
#include <ulocks.h>
#include <task.h>

int i;

main() {
    m_set_procs(2);
    m_fork(func);
    m_kill_procs();
}

void func() {
    int id;

    id=m_get_myid();
    printf("%d - hello world\n",id);
}
```

OUTPUT:

```
1 - hello world
0 - hello world
```

# Process Blocking and Unblocking

---

- Process creation and destruction can be very expensive. On the other hand, keeping processes idle is also expensive
- Alternative - put processes to sleep when not needed and revive them later:

```
main() {  
  
    m_fork(parallel_func);  
    m_kill_procs()  
}  
void parallel_func() {  
  
    (PARALLEL SECTION)  
    m_park_procs();  
    (SERIAL SECTION)  
    m_rele_procs()  
    (PARALLEL SECTION)  
}
```

# Shared Variables

---

- All global variables are shared between processes:

```
foo xxx,yyy,zzz;
```

```
main() {
```

```
    ...
```

```
}
```

```
void parallel_func() {
```

```
    ...
```

```
}
```

- Variables can also be shared using parameter passing

# Example

---

```
float sum,sum0,sum1;

main() {

    m_set_procs(2);
    m_fork(parallel_func);
    m_kill_procs();

    sum=sum0+sum1;
    printf("total sum is %lf\n",sum);
}

void parallel_func() {

    int id;

    id=m_get_myid()
    if (id==0) sum0=1.0+2.0;
    else sum1=3.0+4.0;
}
```

# Contention

---

- Random process scheduling creates problems with respect to modifying shared variables
- The final values of a shared variable can differ from run to run depending on the order in which it was modified
- Root of this contention problem is the simultaneous accessing of a shared variable
- Solution to the contention problem is to provide primitives for the restricted access of shared variables

# Example

---

```
float total_sum;

main() {
    total_sum=0.0;
    m_set_procs(2);
    m_fork(parallel_func);
    m_kill_procs();
    printf("total sum is %f\n",total_sum);
}

void parallel_func() {
    int id;
    float partial_sum;

    id = m_get_myid();
    if (id==0) partial_sum=1.0+2.0;
    else partial_sum=3.0+4.0;

    total_sum=total_sum + partial_sum;
}
```



# Example

---

OUTPUT:

```
total sum is 10.0
(ts = 0;
P0: ts = ts(0) + 3;
P1: ts = ts(3) + 7;)
```

```
total sum is 3.0
(ts = 0;
P1: ts = ts(0) + 7;
P0: ts = ts(0) + 3;)
```

```
total sum is 7.0
(ts = 0;
P0: ts = ts(0) + 3;
P1: ts = ts(0) + 7;)
```

- First output correct because of no contention
- Other outputs wrong because of contention - both processors read the same value of total\_sum (0.0) and one or the other processor updates it last

# Variable Locks

---

- Locks are used to provide exclusive access for the modification of a variable
- A lock variable is to be created for every shared variable that needs exclusive access:
  - It must be shared among all processes that need to use it
  - It can either be in a “locked” or “unlocked” state
  - A locked state indicates ongoing exclusive modification of the variable the lock corresponds to
  - An unlocked state indicates no ongoing modification of the variable

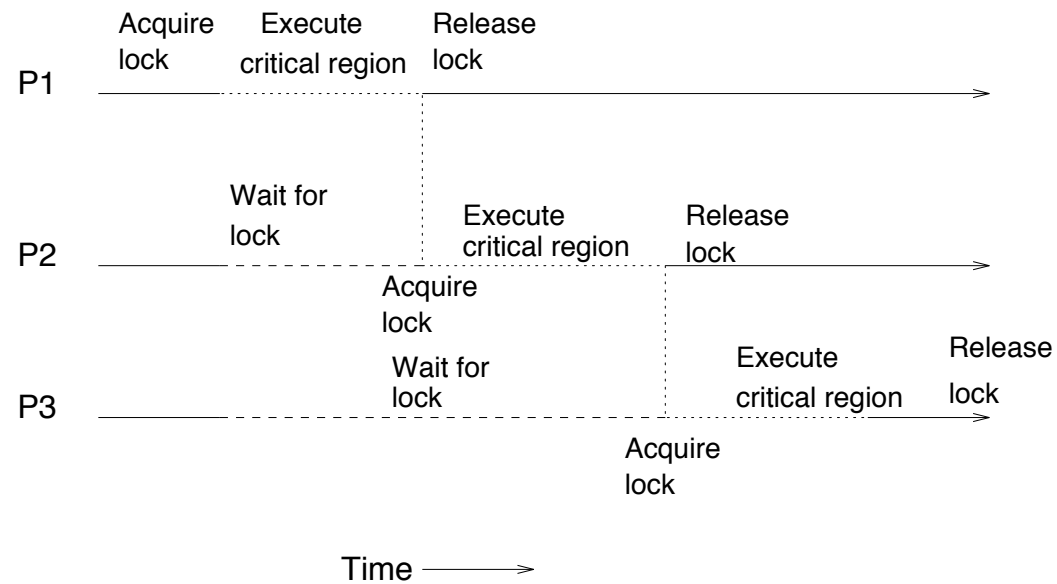
# Variable Locks

---

- The steps to be followed for exclusive access using locks are:
  - Acquire the lock for a variable - first, ensure the lock is in the unlocked state and then switch it to the locked state
  - Modify the variable - read its value and write a new value
  - Release the lock for the variable - switch the lock to the unlocked state so that others may have access
- The use of locks sequentializes execution of program sections and hurts performance

# Illustration of Locks

---



# Example

---

```
float total_sum;

void parallel_func() {
    int id;
    float partial_sum;

    if (id==0) partial_sum=1.0+2.0;
    else partial_sum=3.0+4.0;

    m_lock();
    total_sum+=partial_sum;
    m_unlock();
}
```

# Race Conditions

---

```
float total_sum;

void parallel_func() {
    int id;
    float partial_sum, average;

    if (id==0) partial_sum=1.0+2.0;
    else partial_sum=3.0+4.0;

    m_lock();
    total_sum+=partial_sum;
    m_unlock();

    average=total_sum/4.0;
    printf("%d average is %f\n",id,average);
}
```

OUTPUT:

0 average is 2.5  
1 average is 2.5

1 average is 1.75  
0 average is 2.5

0 average is 0.75  
1 average is 2.5

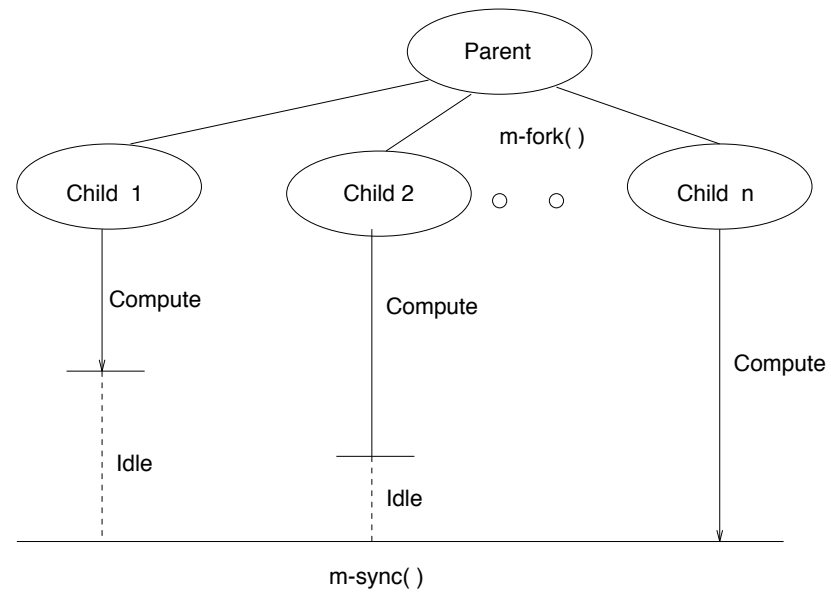
# Barriers

---

- A race condition exists if the results of a parallel program depend on the relative execution speed of processes
- Barriers enable processes to synchronize with each other and help avoid race conditions
- When a process enters a barrier, it waits for all other processes involved to reach the barrier before continuing
- Barriers cause performance degradation and therefore must be used judiciously

# Illustration of Barriers

---





# Barriers

---

```
float total_sum;

void parallel_func() {
    int id;
    float partial_sum, average;

    if (id==0) partial_sum=1.0+2.0;
    else partial_sum=3.0+4.0;

    m_lock();
    total_sum+=partial_sum;
    m_unlock();

    m_barrier();

    average=total_sum/4.0;
    printf("%d average is %f\n",id,average);
}
```

# Process Summary

---

- Process creation/destruction
- Shared variables for inter-process communication
- Locks for exclusive access to shared variables
- Barriers to avoid race conditions

# Process Based Parallel Program Model

---

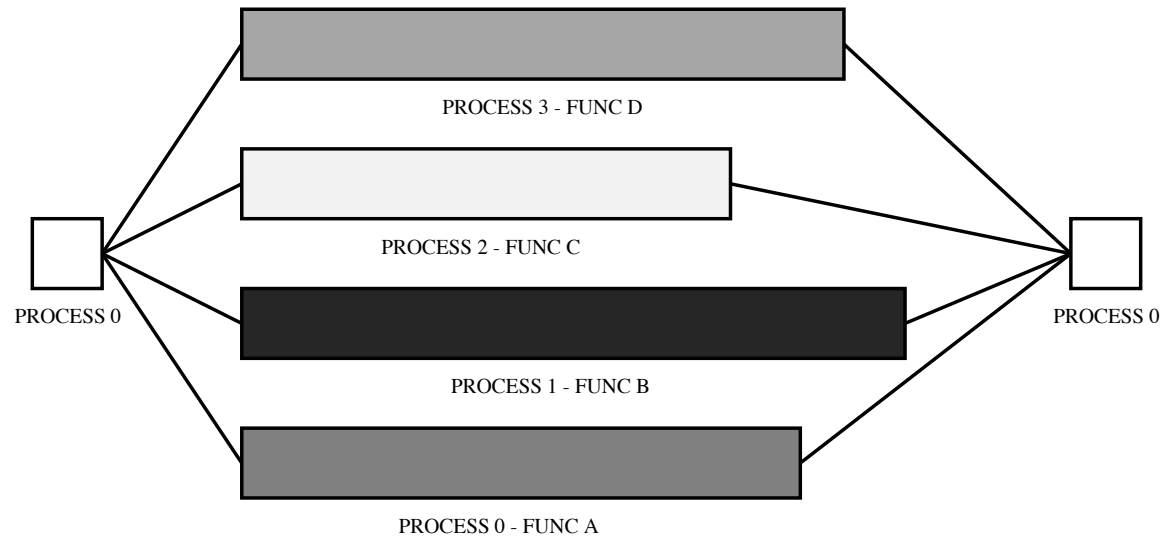
- Processes executing in a function parallel fashion

```
main() {
    m_set_procs(4);
    m_fork(func);
    m_kill_procs();
}

void func() {
    switch (id) {
        case 0 : funcA();
            break;
        case 1 : funcB();
            break;
        case 2 : funcC();
            break;
        case 3 : funcD();
            break;
    }
}
```

# Process Based Parallel Program Model

---



# Process Based Parallel Program Model


---

- Processes executing in a data parallel fashion

```
main() {  
    m_set_procs(4);  
    m_fork(funcA);  
    m_park_procs();  
    ...  
    m_rele_procs();  
    funcB();  
    m_kill_procs();  
}
```

# Process Based Parallel Program Model

---



# Process Based Parallel Program Model

---

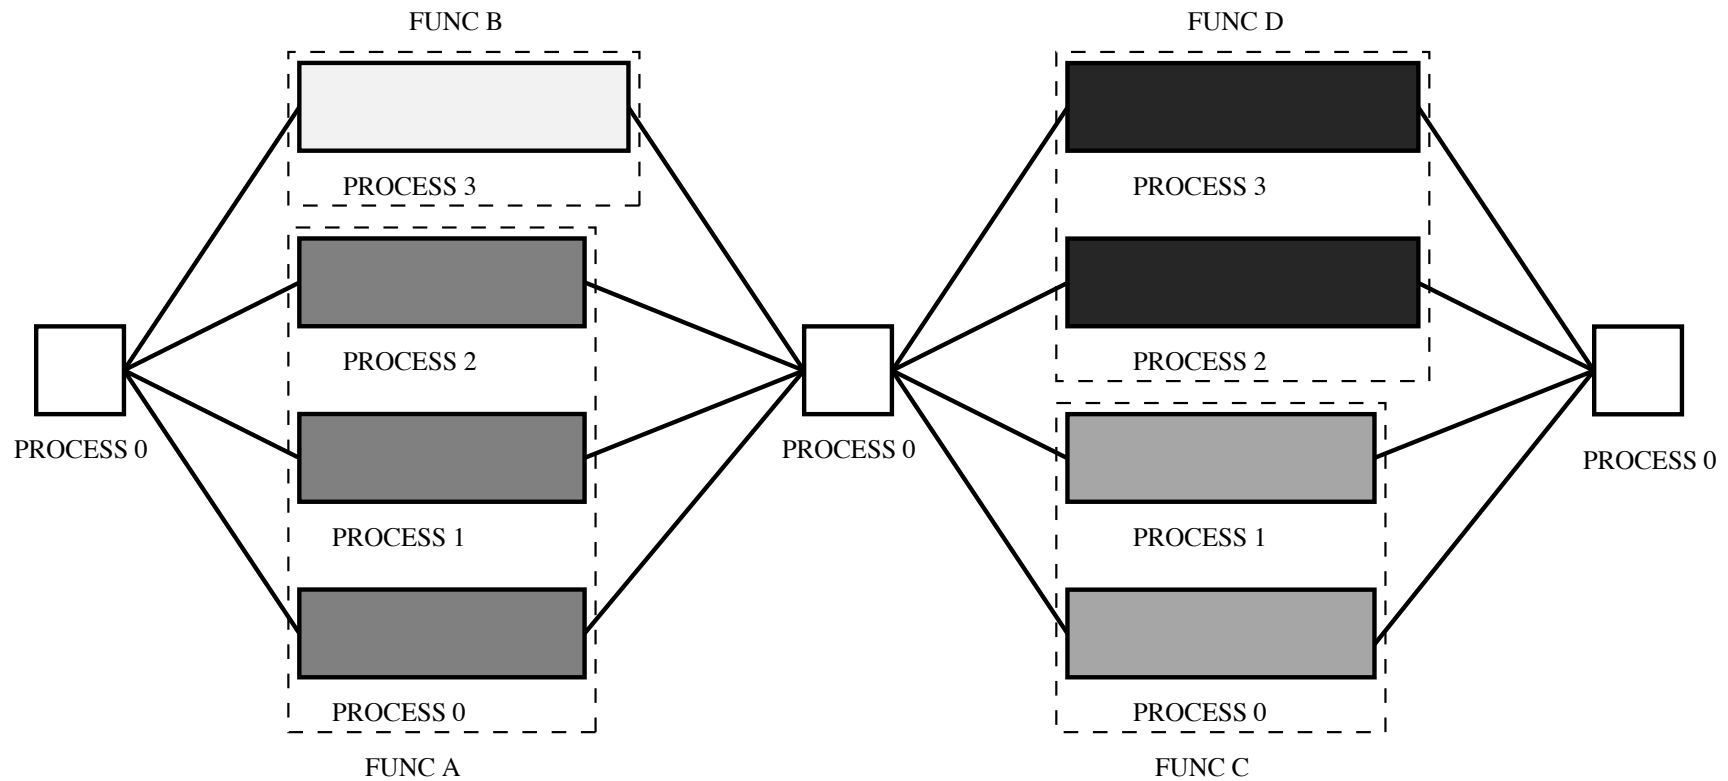
- Processes executing in a function and data parallel fashion

```
main() {  
    m_set_procs(4);  
    m_fork(func1)  
    m_park_procs();  
    ...  
    m_rele_procs();  
    func2;  
    m_kill_procs();  
}
```

```
void func1() {  
    switch (id) {  
        case 0,1,2 : funcA();  
        break;  
        case 3 : funcB();  
        break;  
    }  
}  
  
void func2() {  
    switch (id) {  
        case 0,1 : funcC();  
        break;  
        case 2,3 : funcD();  
        break;  
    }  
}
```

# Process Based Parallel Program Model

---





# Process Based Parallel Program Model

---

- Data parallel programming dominates
- Loops in programs are the source of data parallelism
- Exploitation of parallelism involves sharing work in loops among processes
- Have to use appropriate scheduling techniques for optimal work sharing
- Parallelism in loops not always straightforward to find due to dependence
- Have to perform some transformations to expose parallelism

# Summary

---

- An overview of shared memory parallel programming
- Process model: creation and destruction
- Shared variables
- Synchronization: locks, barriers
- Parallel programming basics
- NEXT CLASS: Shared Memory Parallel Programming - II
- READING: Kumar, Chap. 7