
Electrical Engineering and Computer Science
EECS 358 - INTRODUCTION TO PARALLEL COMPUTING

Lecture 7
Shared Memory Programming - IV

Outline

- Programming with PTHREADS
- Example parallel computation: Matrix vector multiplication
- Programming with parallelization directives
- READING:
 - POSIX PTHREADS Programming Manual.
 - B. Bauer, “Practical Parallel Programming”, chap 7, 12,

Introduction to Multithreaded Programming

- All programmers are used to programming a single thread of control
- A thread is a single sequence of execution steps performed by a program.
- A traditional UNIX process is a program with a single thread of control; it has sole possession of its address space.
- In traditional symmetric multiprocessing operating systems such as IRIX, when a user makes call to *m-fork()*, a new process is created.
- In Linux, you create them in sequence (which adds more cost) by *pthread_create*
- Such processes can share data through the *mmap()* UNIX call, or through the use of *shared* directives.
- Multithreaded libraries allow multiple threads to share an address space.
- Cost of creating a thread is very small
- Hence threads are called “lightweight processes”

Introduction to PTHREADS

- Each parallel computer vendor has its own way of supporting shared memory parallel programming
 - Sequent , Encore, SGI, SUN, HP, IBM, Convex
- IEEE had a POSIX 1003 group that defined a UNIFORM interface to multi-threaded programming
- This is called PTHREADS, and is similar to Solaris Threads from Sun

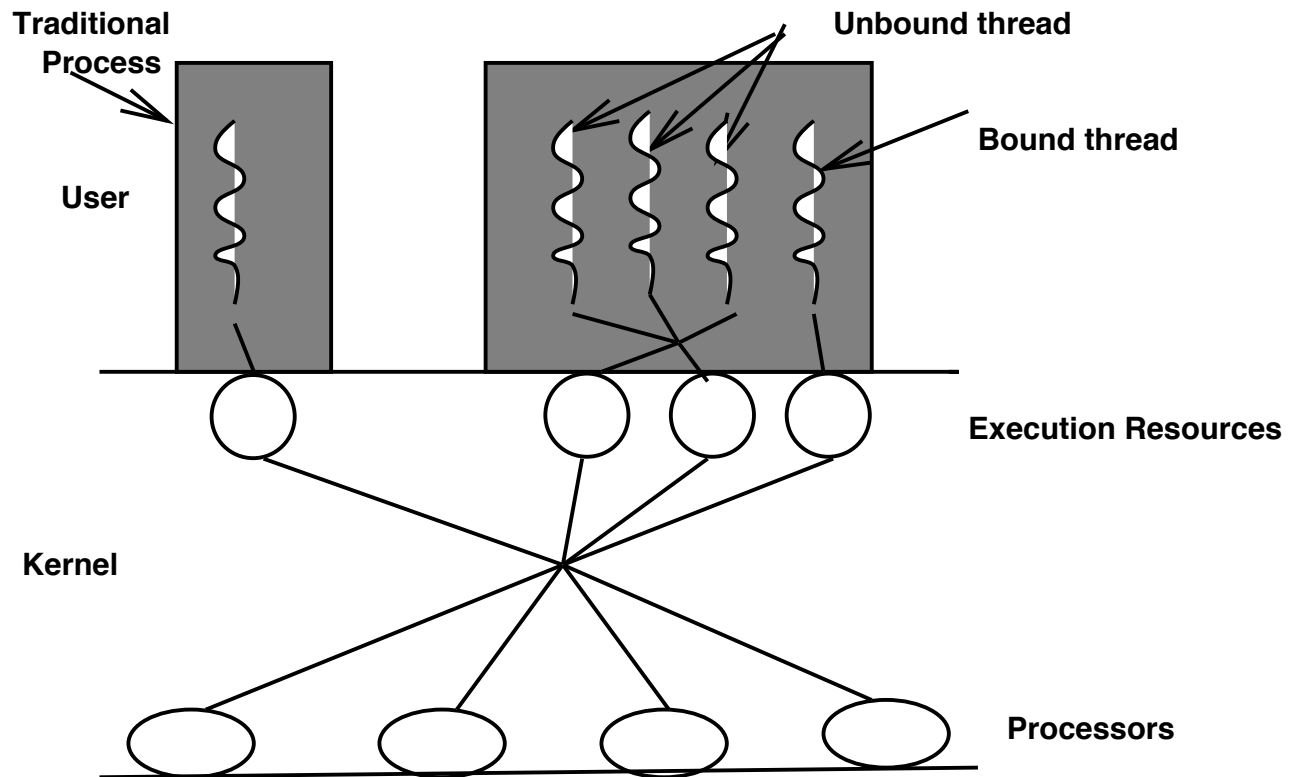
Introduction to PTHREADS Programming

- A multithreaded process is not a thread of control itself. It contains one or more threads of control.
- Some system calls affect the process as a whole
 - Example: if one thread calls the EXIT() call, all the threads in the process are destroyed.
- Most process resources are equally accessible to all the threads in the process.
- All process virtual memory is shared; a change in shared data by one thread is available to the other threads in the process.
- Each thread sees the same open files. If two threads read the same file sequentially, they advance through the file as though one thread were reading it.

Per-Thread Resources

- Each thread has:
 - A thread identifier
 - A set of **registers** (including program counter and stack pointer)
 - The **signal mask**
 - Execution priority
 - **Stack**
 - **Thread-specific data**


General Architecture of Threads Interface



Synchronization Primitives

- PTHREADS provide a variety of synchronization facilities for threads to co-operate in accessing shared resources.
 - Mutual exclusion locks (mutex locks)
 - Condition variables
 - Counting semaphores
 - Multiple readers, single writer locks

Use of Synchronization Variables in Shared Memory



- Synchronization variables S used by various threads in two different processes.

Thread Creation

```
int pthread_create(pthread_t *thread,  
    const pthread_attr_t *attr = NULL,  
    void *(*entry) (void *), void *arg);
```

- The function creates and starts a new thread
- The new thread executes the function specified in the call with the optional arguments.

Other Thread Control

- To determine the thread ID of the calling thread

```
pthread_t pthread_self(void);
```

- To terminate a calling thread

```
void pthread_exit(void *status = NULL);
```

Use of Mutex Locks

- The following function acquires the lock or blocks the calling thread if the lock is already held.

```
int pthread_mutex_lock( pthread_mutex_t *mutex);
```

- The following function enables asynchronous polling of lock. It acquires the lock if free, and does not block (returns error) if the lock is held.

```
int pthread_mutex_trylock( pthread_mutex_t *mutex);
```

- The following function unlocks the mutex lock. If other threads are blocked waiting for the lock, thread at head of queue is unblocked.

```
int pthread_mutex_unlock( pthread_mutex_t *mutex);
```

Example of Use of Locks

```
mutex_t count_mutex;
int count;

increment_count()
{
    pthread_mutex_lock(&count_mutex);
    count = count + 1;
    pthread_mutex_unlock(&count_mutex);
}

int get_count()
{
    int c;

    pthread_mutex_lock(&count_mutex);
    c = count;
    pthread_mutex_unlock(&count_mutex);
    return(c);
}
```

Condition Variables

- Condition Variables are used to wait until a particular condition is true.
- A condition variable must be used in conjunction with a mutex lock.
- The following function blocks the calling thread until the condition is signaled. It atomically releases the associated mutex lock before blocking, and atomically reacquires it before returning.

```
int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);
```

- The following function signals one of the threads blocked in cond_wait().

```
int pthread_cond_signal( pthread_cond_t *cond);
```

Condition Variable Usage

```
pthread_mutex_t count_lock;
pthread_cond_t count_nonzero;
unsigned int count;

decrement_count();
{
    pthread_mutex_lock(&count_lock);
    while (count == 0)
        pthread_cond_wait(&count_nonzero, &count_lock);
    count = count - 1;
    pthread_mutex_unlock(&count_lock);
}

increment_count();
{
    pthread_mutex_lock(&count_lock);
    if (count == 0)
        pthread_cond_signal(&count_nonzero);
    count = count + 1;
    pthread_mutex_unlock(&count_lock);
}
```

Semaphores

- A semaphore is a non-negative integer count that is incremented and decremented.
- It is not as efficient as a mutex lock. They need not be acquired and released by the same thread, so they may be used in asynchronous event notification.
- The following function blocks the thread until the semaphore becomes greater than zero, then decrements it.

```
pthread_sema_wait(pthread_sema_t *sem);
```

- The following function increments the semaphore, potentially unblocking a waiting thread.

```
pthread_sema_post(pthread_sema_t *sem);
```


Use of Semaphores (Producer-Consumer)

```
int rdprt = 0; /* pointers to a circular buffer */
int wrptr = 0; /* of size BUFSIZE */
data_t buf[BUFSIZE];
pthread_sema_t sem;
```

THREAD 1:

```
while (work_to_do) {
    buf[wrptr] = produce();
    wrptr = (wrptr + 1) % BUFSIZE;
    pthread_sema_post(&sem);
}
```

THREAD 2:

```
while (work_to_do) {
    pthread_sema_wait(&sem);
    consume(buf[rdprt]);
    rdprt = (rdprt + 1) % BUFSIZE;
}
```

Example: Matrix Vector Multiplication

```
float a[max][max], b[max], c[max];
int m,n;
pthread_mutex_t vec_lock;

main()

{
    int i,j,nproc;
    void readdata(),matvec(),printdata();

    /* input m and n */

    printf("Enter number of rows and columns");
    scanf("%d %d",&m,&n);
    printf("Enter number of processors\n");
    scanf("%d",&nproc);
    /* read data */
    readdata();
```

Shared Memory Matrix-Vector Multiplication

```
/* set desired level of concurrency */
pthread_setconcurrency(&nprocs);

/* create parallel threads */
for (i=0; i < nprocs; i++)
    pthread_create(NULL, NULL, matvec, (void *)i);

/* print output */
printdata();
}
```

Shared Memory Matrix Vector Multiplication

```
void matvec(void *args)
{
    int i,j,nprocsi,myid;
    float tmp[max];
    for (i=0; i < m; i++) {
        tmp[i] = 0;
        c[i] = 0;
    }
    nprocs = pthread_getconcurrency();
    myid = (int)args;
    for (j = myid; j < n; j = j + nprocs) {
        for (i=0; i < m; i++)
            tmp[i] = tmp[i] + a[i][j] * b[j];
    }
    pthread_lock(&vec_lock);
    for (i=0; i < m; i++)
        c[i] = c[i] + tmp[i];
    pthread_unlock(&vec_lock);
}
```

How to Program with PTHREADS

- To program a parallel application with PTHREADS, need to add this statement to the source file:

```
#include <pthread.h>
```

- Compile a C program using

```
% cc input.c -lpthread
```

Parallelization Directives

- Until now, we have created parallel program using:
 - Creation and management of processes
 - Implementation of scheduling scheme
- Observations:
 - Procedure tedious and routine, making changes difficult
- Directives are used to inform the compiler about parallelism
- We will look at OpenMP (Open Multi-Processing) parallelization directives

Parallelization Directives

- Directive used before a loop tells compiler how to parallelize the code in the loop:

```
#pragma parallel
[options]
#pragma pfor iterate(i=0;n;1)
[options]
for (i=0;i<n;i++) {
    ...
}
```

- Linux:

```
#pragma omp parallel [options]
{
    #pragma omp for [options]
    for (...)
}
```

#pragma Options

- **private(list):**
 - Variables in the list are local to each process
 - The value of the private variables is undefined beyond the loop
 - All loop indices are private by default, all other variables are shared by default
- **lastprivate(list):**
 - Variables in the list are local to each process
 - The values of these private variables beyond the loop is guaranteed to be that obtained in the last iteration

#pragma Options

- **share(list):**
 - Variables in the list are shared among processes
 - All variables except the loop indices default to shared if not explicitly declared otherwise
- **if(expression):**
 - Provides a run-time selection between parallel and serial execution of the loop
 - If the expression evaluates to 1, the loop executes in parallel

#pragma parallel Options

- **byvalue(list):**
 - Variables in list are treated as read only variables
 - The compiler makes copies of these variables for each process
- **numthreads(num):**
 - Sets the number of processes on which the loop is executed to num
 - Default value is the number of processors available

#pragma pfor Options

- **iterate(var=start;num_iters;step):**
 - Assigns the index variable and sets the initial value, total number of iterations and stepsize
- **mp_schedtype=type:**
 - Sets the scheduling scheme for the loop
 - Types can be:
 - * *simple* - static block
 - * *interleaved* - static interleaved
 - * *dynamic* - self scheduling
 - * *gss* - guided self scheduling
- **chunk=num:**
 - Sets the chunksize for dynamic and interleaved scheduling schemes

Local Variables

```
for (i=0;i<n;i++) {  
    x=a[i]*3;  
    b[i]=x*b[i];  
}
```

```
#pragma parallel private (x,i)  
#pragma shared(a,b)  
#pragma byvalue(n)  
{  
    #pragma pfor iterate(i=0;n;1)  
    for (i=0;i<n;i++) {  
        x=a[i]*3;  
        b[i]=x*b[i];  
    }  
}
```

Local Variables

```
for (i=0;i<n;i++) {
    x=a[i]**3;
    b[i]=x*b[i];
}
val=foo(x);

#pragma parallel private (local_x,i)
#pragma shared(a,b,x)
#pragma byvalue(n)
{
    #pragma pfor iterate(i=0;n;1)
    for (i=0;i<n;i++) {
        local_x=a[i]**3;
        b[i]=local_x*b[i];
        if (i==n-1) x=local_x;
    }
}
val=foo(x);
```

Loop-Carried Values

```
indx=0;
for (i=0;i<n;i++) {
    indx+=i;
    a[i]=b[indx];
}

#pragma parallel private (indx,i)
#pragma shared(a,b)
{
    #pragma pfor iterate(i=0;n;1)
    #pragma schedtype(simple)

    for (i=0;i<n;i++) {
        indx=(i*(i+1))/2;
        a[i]=b[indx];
    }
}
```

Sum Reduction

```
total=0.0
for(i=0;i<n;i++) {
    total+=a[i];
}
```

```
total=0.0;
#pragma parallel private (sub_total,i)
#pragma shared(a,total)
{
    sub_total=0.0;
    #pragma pfor iterate(i=0;n;1)
    for (i=0;i<n;i++)
    {
        sub_total+=a[i];
    }
    #pragma critical
    {
        total+=sub_total;
    }
}
```

Summary

- Programming with PTHREADS
- Example parallel computation: Matrix vector multiplication
- Programming with parallelization directives
- NEXT LECTURE: Distributed Memory Parallel Architectures.
- READING: V. Kumar et al, "Introduction to Parallel Computing," Chapter 6.