

---

**Electrical Engineering and Computer Science**  
**EECS 358 - INTRODUCTION TO PARALLEL COMPUTING**

**Lecture 12**

**Dist. Mem. Message Passing Programming - III**

---

# Outline

---

- Intermediate MPI concepts
- Solution of Poisson Problem using Jacobi Method
- Use of Topologies
- Use of Decompositions
- Use of Data Exchange Routines
- Use of Derived Datatypes
- READING: Foster, Chapter 8

# The Poisson Problem

---

- The Poisson Problem is a simple partial differential equation (PDE)

$$\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = f(x, y) \quad (1)$$

given  $u(x, y) = g(x, y)$  on the boundary.

- We may "solve" the Poisson equation numerically over a region by discretizing it in the  $x$  and  $y$  directions to obtain a grid of points and then computing the approximate solution values at these points.
- Assuming that the distances between neighboring points in the grid in the  $x$  and  $y$  directions is  $h$

# The Poisson Problem

---

- We may replace the partial derivatives by numerical approximations involving first-order finite differences to get


$$u(i-1, j) + u(i, j+1) + u(i, j-1) + u(i+1, j) - 4u(i, j) = f(i, j) \quad (2)$$

- We can then write a Jacobi iteration as

$$u^{k+1}(i, j) = 1/4(u^k(i-1, j) + u^k(i, j+1) + u^k(i, j-1) + u^k(i+1, j) - h^2 f(i, j)) \quad (3)$$

# Poisson Problem

---



# Finite Difference Algorithm

---

```
int finite_difference()
{
    int i, j, n;
    double u[n,n], unew[n,n];
    for (k = 0; k < ITERS; k++) {
        for (i = 0; i < n; i++) {
            for (j = 0; j < n; j++) {
                unew[i,j] = 0.25 * (u[i-1,j] + u[i+1,j]
                                   + u[i,j-1] + u[i,j+1]) - f[i,j];
            }
        }
        diffmax = 0.0;
        for (i = 1; i < n; i++) {
            for (j = 1; j < n; j++) {
                diff = abs(unew[i,j] - u[i,j]);
                if (diff > diffmax) diffmax = diff;
                u[i,j] = unew[i,j];
            }
        }
    }
}
```


# Basic Sweep Routine

---

```
int sweep()
{
int i, j, n;
double u[n,n], unew[n,n];
    for (i = 0; j < n; j++) {
        for (j = 0; i < n; i++)
            unew[i,j] = 0.25 * (u[i-1,j] + u[i+1,j]
                                + u[i,j-1] + u[i, j+1]) - f[i,j];
    }
```

# One-dimensional Decomposition of Domain

---





# Jacobi Iteration for Slice of Domain

---

```
int i, j, n, m;
    /* assume P processors, partition array by rows */
double u[n/P,n], unew[n/P,n];

for (i = 0; i < n/P; i++)
    for (j = 0; j < n; j++) {
        unew[i,j] = 0.25 * (u[i-1,j] + u[i+1,j]
                           + u[i,j-1] + u[i, j+1]) - f[i,j];
    }
```

## Problem with Previous Code


---

- The loop will require elements such as  $u[i-1,j]$  or  $u[i+1,j]$  from a different process
- We will discuss how to get this data from other processes
- Define an overlap region to hold data from other regions

```
double u[n/P+2, n]
```

# Computational Domain with Overlap Regions

---



# Defining Topologies

---


- Our next task is to define how to assign processes to each part of the decomposed domain
- MPI lets user specify various *application topologies*
- The routine `MPI_Cart_create()` creates a Cartesian decomposition of the processes, with the number of dimensions given by the *ndim* argument

```
dims[0] = 4
dims[1] = 3
periods[0] = 0
periods[1] = 0
reorder = 1
ndim = 2
MPI_Cart_create(MPI_COMM_WORLD, ndim, *dims,
                *periods, reorder, comm2d)
```

- This creates a new communicator “comm2d”
- The “periods” argument specifies if connection is with wraparound, i.e. if the processes at end are connected together.

# Illustration of Cartesian Topology

---



# Finding Neighbors of a Topology

---

- To determine the coordinates of a calling process

```
MPI_Cart_get(comm1d, 2, *dims, *periods, *coords);  
printf('( ', coords[0], ', ' coords[1], ')');
```

- To determine rank

```
MPI_Cart_rank(comm1d, *coord, *myrank);  
printf(myrank);
```

# Using Topologies


---

- Performing shift operations

```
MPI_Cart_shift(comm1d, direction, shift, *src, *dest);
```

# Data Exchange Operations

---





# Data Exchange Routine

---

```
int exchngr1(a, nx, comm1d, nbrbottom, nbrtop)
{
    int nx; double a(n/P+2,n)
    int comm1d, nbrbottom, nbrtop
    int status(MPI_STATUS_SIZE), ierr

    if (nbrtop != MPI_UNDEFINED) {
        MPI_Send(&a(1,0), nx, MPI_DOUBLE,
                 nbrtop, 0, comm1d);
        MPI_Recv(&a(0,0), nx, MPI_DOUBLE,
                 nbrtop, 0, comm1d);
    }

    if (nbrbottom != MPI_UNDEFINED) {
        MPI_Send(&a(n/P,0), nx, MPI_DOUBLE,
                 nbrbottom, 1, comm1d);
        MPI_Recv(&a(n/P+1,0), nx, MPI_DOUBLE,
                 nbrbottom, 1, comm1d);
    }
}
```

# Basic Jacobi Sweep Routine

---

```
int sweep1d( &a, &f, nx, &b)
int nx
double a(n/P+2,n), f(n/P+2,n), b(n/P+2,n)

int i, j
double h

h = 1.0/(nx+1);

for (i = 1; i < (n/P+1); i++)
    for (j = 0; j < n; j++)
        b[i,j] = 0.25 * (a[i-1,j] + a[i+1,j]
            + a[i,j-1] + a[i, j+1]) - h * h * f[i,j])
```

# Overall Jacobi Implementation

---

```
main()
    #define maxn 128
    double a(maxn,maxn), b(maxn,maxn), f(maxn,maxn);
    int nx, ny;
    int myid, numprocs;
    int commld, nbrbottom, nbrtop, it;

    MPI_Init();
    MPI_Comm_rank( MPI_COMM_WORLD, myid);
    MPI_Comm_size( MPI_COMM_WORLD, numprocs);
    if (myid == 0) nx = 100;
    MPI_Bcast(nx,1,MPI_INT,0,MPI_COMM_WORLD);
    ny      = nx;
```

## Jacobi Implementation-contd

---

```
/* Create a one dim cartesian mapping
    MPI_Cart_Create( MPI_COMM_WORLD, 1, numprocs, 0,1,comm1d);
/* Get my position in this communicator, and my neighbors
    MPI_Comm_Rank(comm1d, myid);
    MPI_Cart_shift( comm1d, 0,1, &nbrbottom, &nbrtop);
/* Initialize the right-hand-side (f) and the initial solution guess (a)
    onedinit( a, b, f, nx);
```

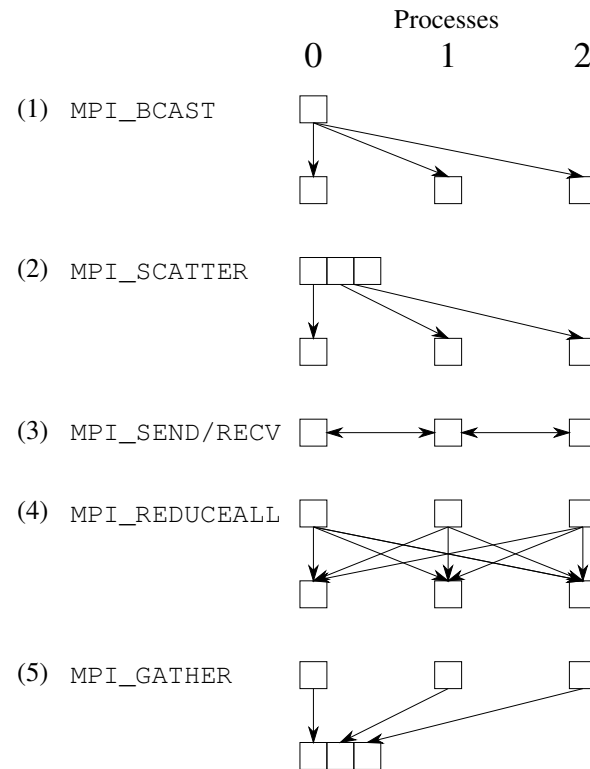
## Jacobi Implementation-contd

---

```
MPI_Barrier( MPI_COMM_WORLD);
for (it=1; it < 100; it++) {
    exchng1( a, nx, comm1d, nbrbottom, nbrtop );
    sweep1d( a, f, nx, b );
    dwork = diff( a, b, nx, s, e );
    MPI_Allreduce( dwork, diffnorm, 1,MPI_DOUBLE,MPI_MAX,comm1d);
    if (diffnorm < 1.0e-5) break;
    exchng1( b, nx, comm1d, nbrbottom, nbrtop );
    sweep1d( b, f, nx, a );
    dwork = diff( a, b, nx, s, e );
    MPI_Allreduce( dwork, diffnorm, 1,MPI_DOUBLE,MPI_MAX,comm1d);
    if (diffnorm < 1.0e-5) break;
}
MPI_Finalize();
```

# Overall Communication Pattern of Jacobi

---



# Use of MPI DataTypes

---

- One of MPI's novel features is the use of a datatype associated with every message.
- Specify the length of a message as a given count of occurrences of a given datatype.
- MPI has basic datatypes such as MPI\_INT, MPI\_FLOAT, etc.
- In MPI a datatype is an object that specifies a sequence of basic datatypes and displacements in bytes for each datatype.

# MPI Datatypes in C

---

MPI Datatype	C Datatype
MPI_BYTE	
MPI_CHAR	signed char
MPI_DOUBLE	double
MPI_FLOAT	float
MPI_INT	int
MPI_LONG	long
MPI_LONG_DOUBLE	long double
MPI_PACKED	
MPI_SHORT	short
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long
MPI_UNSIGNED_SHORT	unsigned short



# Typemaps in MPI

---

- We represent a datatype as a sequence of pairs of basic types and displacements
- Call this sequence a *typemap*
- Typemap =  $(type_0, disp_0), \dots, (type_{n-1}, disp_{n-1})$
- Example: (int,0), (char, 4) is a typemap describing a datatype which is an *integer* of 4 bytes, and then a *character* of 1 byte.
- Strictly, the total size is 5 bytes, but if we assume 4-byte alignment, then using padding, the extent of the previous datatype is 8 bytes.
- To add another *integer*, we have to specify (int,0), (char, 4), (int,8)

# Derived Datatypes

---

- The typemap is a general way of describing an arbitrary datatype.
- However, need a convenient way to represent datatypes with large number of entries.
- CONTIGUOUS: This produces a new datatype by making *count* copies of an existing one.
- VECTOR: This creates a new datatype but allows regular gaps in the displacements
- INDEXED: In this datatype, an array of displacements of the input datatype is provided.

## Example of CONTIGUOUS Datatype

---

- Assume an original datatype *oldtype* has typemap (int, 0), (double, 4), then  
`MPI_Type_contiguous(2, oldtype, & newtype);`
- Creates a datatype *newtype* with a typemap:  
(int, 0), (double, 4),(int, 12), (double, 16) with an extent of 24 bytes.
- To actually send such a data use the sequence of calls:

```
MPI_Type_contiguous(count, datatype, &newtype);  
MPI_Type_commit(&newtype);  
MPI_Send(buffer, 1, newtype, dest, tag, comm);  
MPI_Type_free (&newtype);
```

## Example of VECTOR Datatype

---

```
MPI_Type_Vector(count, blocklength, stride,  
                oldtype, stridetype);  
MPI_Type_Commit(stridetype);  
MPI_Send();
```

- Specifically, for the Jacobi 2-D decomposition

```
MPI_Type_Vector(ey+1 - (sy-1), 1, ex+1 - (sx-1),  
                MPI_DOUBLE, stridetype)  
MPI_Type_Commit(stridetype);  
MPI_SEND();
```

## Example of VECTOR Datatype

---


29	30	31	32	33	34	35
22	23	24	25	26	27	28
15	16	17	18	19	20	21
8	9	10	11	12	13	14
1	2	3	4	5	6	7

`MPI_Type_Vector(5, 1, 7, MPI_DOUBLE, newtype);`

# More Examples of Data Types

---

- Consider program to communicate the north and south rows and east and west columns of a 4 X 6 C array for a finite difference code



# Finite Difference Code Example

---

```
integer coltype, rowtype, comm, ierr;
/* The derived type rowtype is 6 contiguous reals. */
MPI_Type_Contiguous(6, MPI_FLOAT, rowtype);
MPI_Type_Commit(rowtype) ;
/* The derived type coltype is 4 reals, located 6 apart. */
MPI_Type_Vector(4, 1, 6, MPI_FLOAT, coltype, ierr);
MPI_Type_Commit (coltype, ierr);
...
MPI_Send(&array[0,0], 1, coltype, west, 0, comm);
MPI_Send(&array[0,5], 1, coltype, east, 0, comm);
MPI_Send(&array[0,0], 1, rowtype, north, 0, comm);
MPI_Send(&array[3,0], 1, rowtype, south, 0, comm);
...
MPI_Type_Free(rowtype, ierr)
MPI_Type_Free(coltype, ierr)
```

# Summary

---

- Intermediate MPI concepts
- Solution of Poisson Problem using Jacobi Method
- Use of Topologies
- Use of Decompositions
- Use of Data Exchange Routines
- Use of Derived Datatypes