

---

**Electrical Engineering and Computer Science**  
**EECS 358 - INTRODUCTION TO PARALLEL COMPUTING**

**Lecture 15**  
**Dense Matrix Algorithms**

---

# Outline

---

- Design of parallel algorithms
- Dense matrix representations
- Matrix transpose algorithms
- Matrix vector multiplication
- Matrix matrix multiplication
- Linear system of equation solvers
- READING: Chapter 5, V. Kumar, “Introduction to Parallel Computing”

# Design of Parallel Algorithms

---

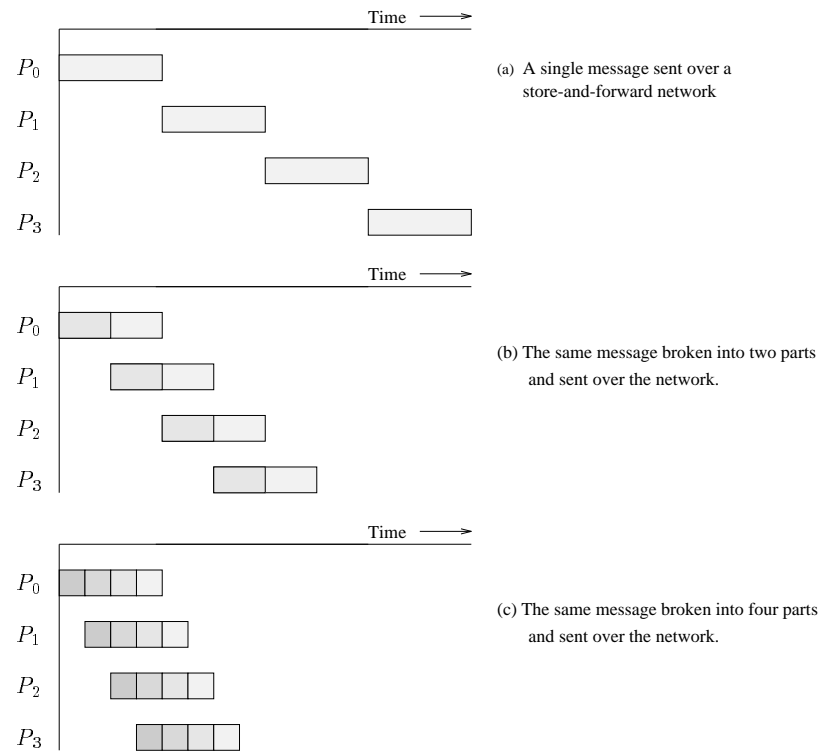
- Assume performance model of cost of communication
  - Startup time for messages ( $t_s$ )
  - Per-hop time ( $t_h$ )
  - Per-word transfer time ( $t_w$ )
- Assume performance model of computation
  - Cost of multiplication or addition or any operation ( $t_c$ )

# Store and Forward vs Cut Through Routing

---

- Cost of store and forward routing of message of size  $m$  via  $l$  hops,  $t_{comm} = t_s + (m.t_w + t_h)l$
- In current parallel computers,  $t_h$  is small, hence approximate as  $t_{comm} = t_s + m.t_w.l$
- Cost of cut through routing of message of size  $m$  via  $l$  hops,  $t_{comm} = t_s + l.t_h + m.t_w$

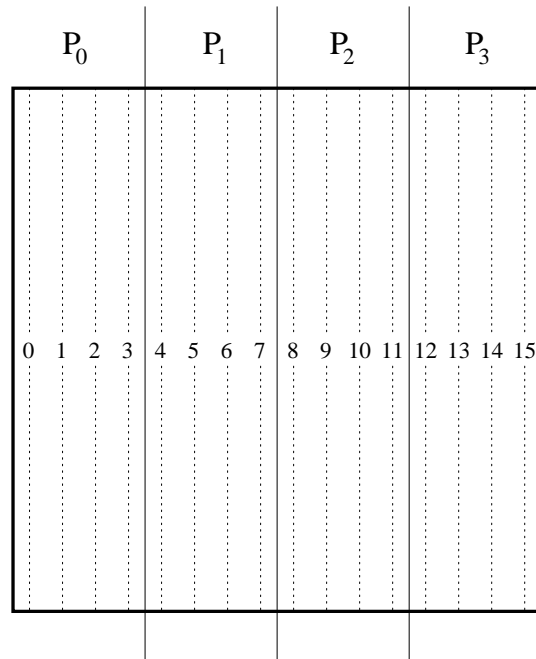
# Store and Forward vs Cut Through Routing



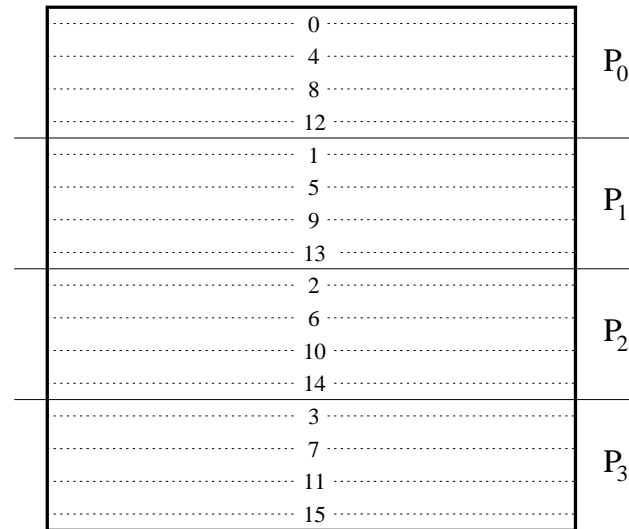
# Dense Matrix Mappings

---

- Striped partitioning (one dimensional, row or column, block or cyclic)



(a) Columnwise block striping



(b) Rowwise cyclic striping

# Dense Matrix Mappings-Continued

---

- Checkerboard partitioning (two dimensional, blocked or cyclic)

(0,0) (0,1)	(0,2) (0,3)	(0,4) (0,5)	(0,6) (0,7)
$P_0$	$P_1$	$P_2$	$P_3$
(1,0) (1,1)	(1,2) (1,3)	(1,4) (1,5)	(1,6) (1,7)
(2,0) (2,1)	(2,2) (2,3)	(2,4) (2,5)	(2,6) (2,7)
$P_4$	$P_5$	$P_6$	$P_7$
(3,0) (3,1)	(3,2) (3,3)	(3,4) (3,5)	(3,6) (3,7)
(4,0) (4,1)	(4,2) (4,3)	(4,4) (4,5)	(4,6) (4,7)
$P_8$	$P_9$	$P_{10}$	$P_{11}$
(5,0) (5,1)	(5,2) (5,3)	(5,4) (5,5)	(5,6) (5,7)
(6,0) (6,1)	(6,2) (6,3)	(6,4) (6,5)	(6,6) (6,7)
$P_{12}$	$P_{13}$	$P_{14}$	$P_{15}$
(7,0) (7,1)	(7,2) (7,3)	(7,4) (7,5)	(7,6) (7,7)

(a) Block-checkerboard partitioning

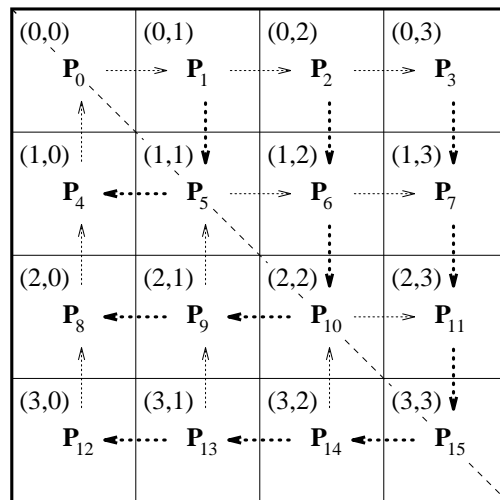
(0,0) (0,4)	(0,1) (0,5)	(0,2) (0,6)	(0,3) (0,7)
$P_0$	$P_1$	$P_2$	$P_3$
(4,0) (4,4)	(4,1) (4,5)	(4,2) (4,6)	(4,3) (4,7)
(1,0) (1,4)	(1,1) (1,5)	(1,2) (1,6)	(1,3) (1,7)
$P_4$	$P_5$	$P_6$	$P_7$
(5,0) (5,4)	(5,1) (5,5)	(5,2) (5,6)	(5,3) (5,7)
(2,0) (2,4)	(2,1) (2,5)	(2,2) (2,6)	(2,3) (2,7)
$P_8$	$P_9$	$P_{10}$	$P_{11}$
(6,0) (6,4)	(6,1) (6,5)	(6,2) (6,6)	(6,3) (6,7)
(3,0) (3,4)	(3,1) (3,5)	(3,2) (3,6)	(3,3) (3,7)
$P_{12}$	$P_{13}$	$P_{14}$	$P_{15}$
(7,0) (7,4)	(7,1) (7,5)	(7,2) (7,6)	(7,3) (7,7)

(b) Cyclic-checkerboard partitioning

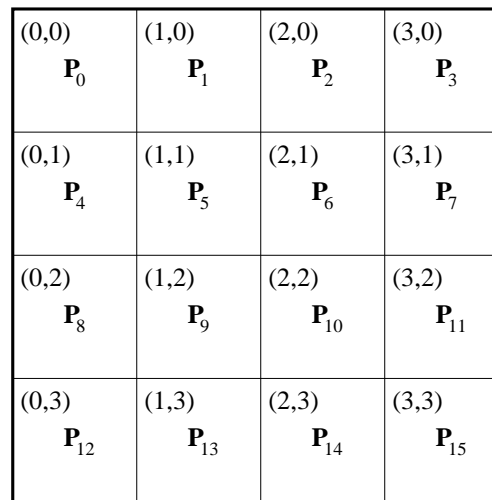
# Matrix Transposition

---

- Need to perform  $A[i,j] = A[j,i]$  for all  $i,j$
- Checkerboard partitioning on a Mesh of processors
- Case 1: number of processors equal to array elements



(a) Communication steps

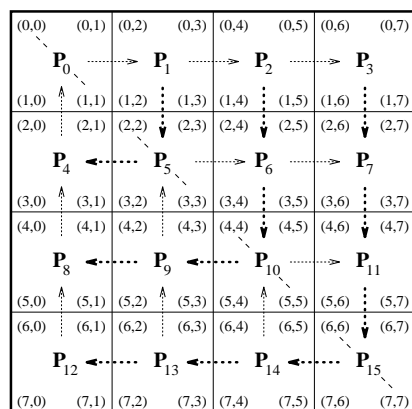


(b) Final configuration

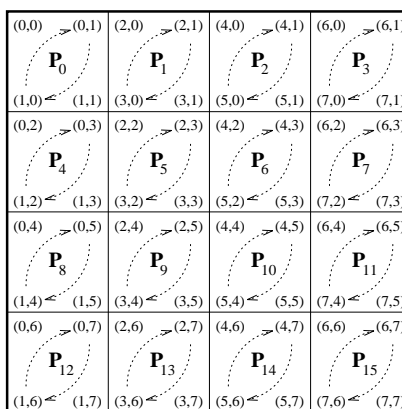


# Matrix Transposition on a Mesh

- Case 2: number of processors less than array elements
- First phase: communicate blocks, second phase, local transpose on blocks
- $T_P = \frac{n^2}{2p} + 2t_s\sqrt{p} + 2t_w\frac{n^2}{\sqrt{p}}$

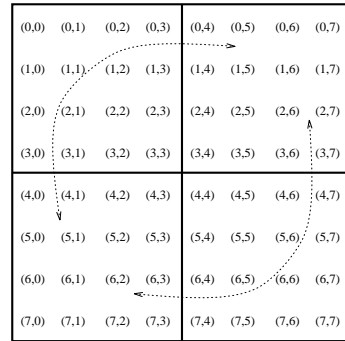


(a) Communication steps

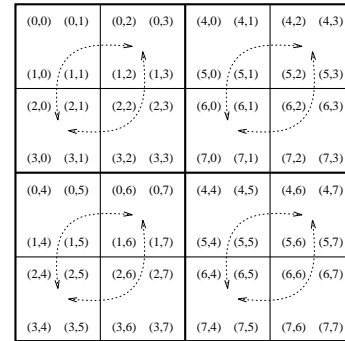


(b) Local rearrangement

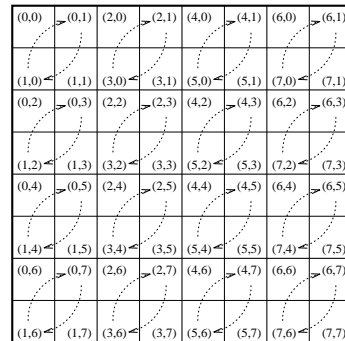
# Recursive Transposition Algorithm



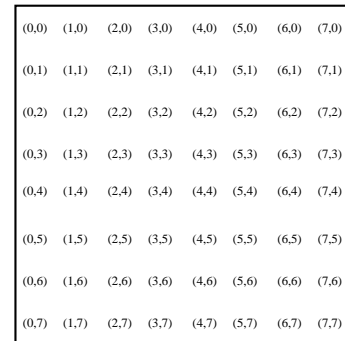
(a) Division of the matrix into four blocks and exchange of top-right and bottom-left blocks



(b) Division of each block into four subblocks and exchange of top-right and bottom-left subblocks



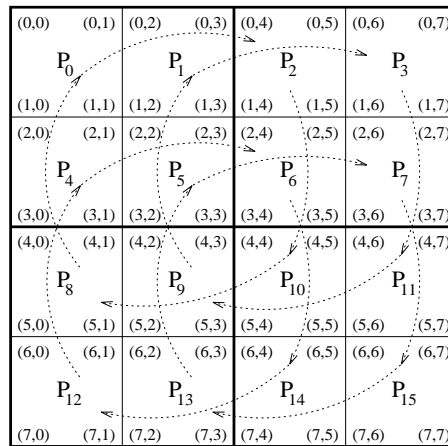
(c) Last subdivision and transposition



(d) Final configuration

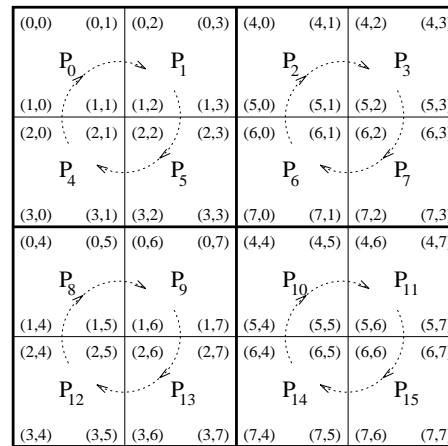
# Matrix Transpose on the Hypercube

- $T_P = \frac{n^2}{2p} + (t_s + t_w \frac{n^2}{p}) \log(p)$



(a)

Division of the matrix into four blocks and exchange of top-right and bottom-left blocks

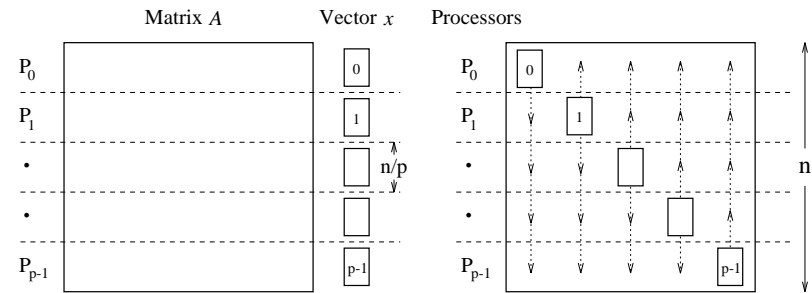


(b)

Division of each block into four subblocks and exchange of top-right and bottom-left subblocks

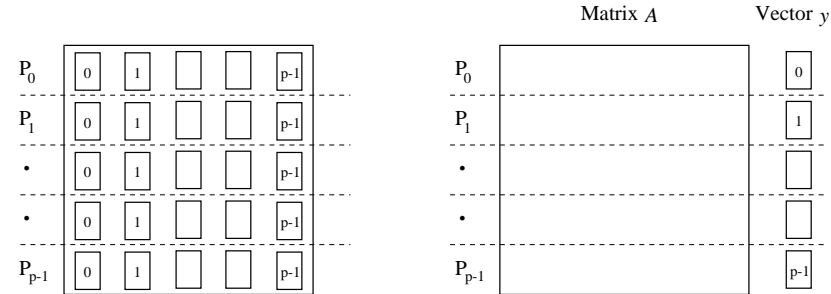
# Matrix Vector Multiplication ( $Ax = y$ )

- Assume row-wise striping, one row per processor



(a) Initial partitioning of the matrix and the starting vector  $x$

(b) Distribution of the full vector among all the processors by all-to-all broadcast



(c) Entire vector distributed to each processor after the broadcast

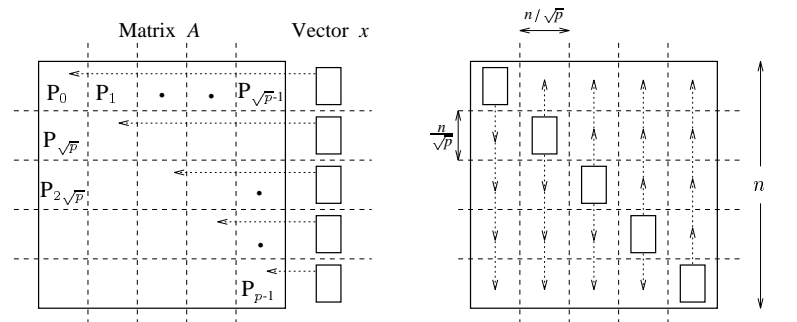
(d) Final distribution of the matrix and the result vector  $y$

# Runtime Analysis

---

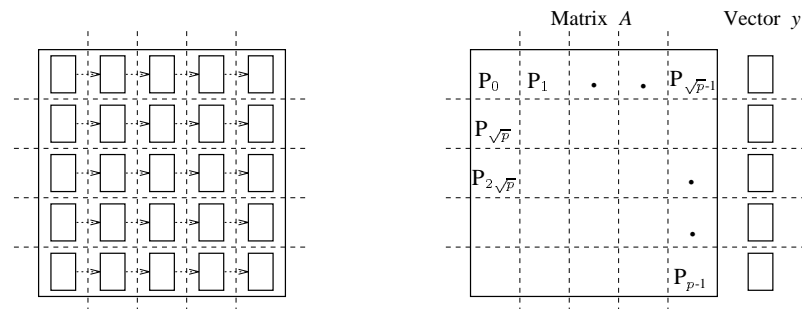
- Assume number of processors less than number of rows
- Runtime on hypercube  $T_P = \frac{n^2}{p} + t_s \log(p) + t_w \cdot (n - \frac{n}{p})$
- Runtime on torus  $T_P = \frac{n^2}{p} + 2 * t_s (\sqrt{p} - 1) + t_w (n - \frac{n}{p})$

# Matrix Vector Using Checkerboard Partitioning



(a) Initial data distribution and communication steps to align the vector along the diagonal

(b) One-to-all broadcast of portions of the vector along processor columns



(c) Single-node accumulation of partial results

(d) Final distribution of the result vector

# Matrix Vector Analysis

---

- Runtime on a mesh with cut-through routing  $T_P = \frac{n^2}{p} + t_s \log(p) + t_w \frac{n}{\sqrt{p}} \log(p) + 2 * t_h(\sqrt{p} - 1)$
- Based on the observation that the broadcast is performed on a ring
- Checkerboard (two-dim blocked) partitioning is faster than striped (one-dim blocked) for both mesh and hypercube architectures

# Matrix Matrix Multiplication ( $C = A.B$ )

---

- Assume  $n \times n$  matrices, serial algorithm takes time  $= n^3$ .
- Consider blocked matrix multiplication algorithm

---

```
1.  procedure BLOCK_MAT_MULT ( $A, B, C$ )
2.  begin
3.      for  $i := 0$  to  $q - 1$  do
4.          for  $j := 0$  to  $q - 1$  do
5.              begin
6.                  Initialize all elements of  $C_{i,j}$  to zero;
7.                  for  $k := 0$  to  $q - 1$  do
8.                       $C_{i,j} := C_{i,j} + A_{i,k} \times B_{k,j}$ ;
9.                  endfor;
10. end BLOCK_MAT_MULT
```

---

**Program 5.3** The block matrix multiplication algorithm for  $n \times n$  matrices with a block size of  $(n/q) \times (n/q)$ .

Copyright (r) 1994 Benjamin/Cummings Publishing Co.



# Simple Parallel Matrix Multiplication

---

- Consider two  $n \times n$  matrices  $A$  and  $B$  partitioned into  $p$  blocks arranged as  $\sqrt{p} \times \sqrt{p}$  logical mesh of processors
- Processor  $P_{i,j}$  initially stores  $A_{i,j}$  and  $B_{i,j}$  computes block  $C_{i,j}$  of result matrix.
- In subsequent stages,  $C_{i,j}$  computation requires all submatrices  $A_{i,j}$  and  $B_{k,j}$  for all  $0 \leq k \leq \sqrt{p}$ .
- An all-to-all broadcast is performed of matrix  $A$ 's blocks for each row of the processors, and an all-to-all broadcast of matrix  $B$ 's blocks in each column.
- Parallel run time on hypercube  $T_P = \frac{n^3}{p} + t_s \log(p) + 2 * t_w \frac{n^2}{p} (\sqrt{p} - 1)$
- Parallel run time on torus  $T_P = \frac{n^3}{p} + 2 * t_s (\sqrt{p} - 1) + 2 * t_w \frac{n^2}{p} (\sqrt{p} - 1)$

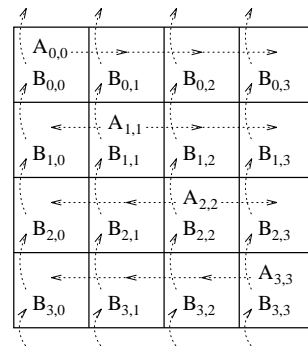
# Memory Efficient Matrix Multiplication

---

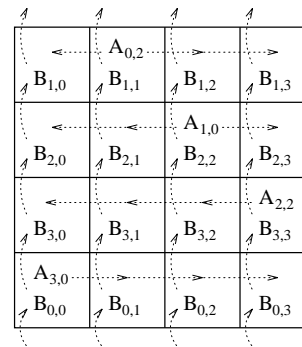
- Partition matrices  $A$  and  $B$  among  $p$  processors so each processor stores  $(n/\sqrt{p}$  by  $(n/\sqrt{p}$  blocks of each matrix.
- Algorithm performs  $\sqrt{p}$  iterations of the following steps:
- Broadcast the selected block of  $A$  among the  $\sqrt{p}$  processors of the row in which the block lies (initially start with block  $A_{i,i}$ ).
- Multiply the block of  $A$  received with the resident block of  $B$ .
- Send the block of  $B$  to processor directly above it (with wraparound), and received new block of  $B$  from processor below it.
- Select the block of  $A$  for the next row broadcast.
- Parallel runtime of algorithm on the hypercube

$$T_P = \frac{n^3}{p} + t_s \sqrt{p} \log(\sqrt{p}) + t_w \frac{n^2}{\sqrt{p}} \log(\sqrt{p})$$

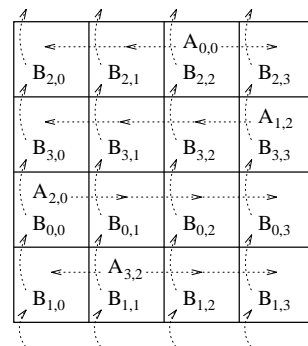
# Memory Efficient Matrix Multiplication



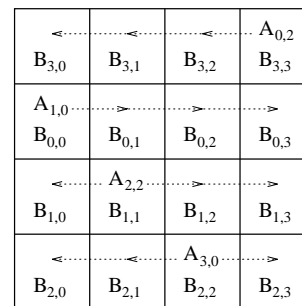
(a)



(b)



(c)



(d)

# Solving System of Linear Equations

---

- Given system of linear equations  $A.x = b$ , solve for vector  $x$  in two stages.
- Reduce system of equations to a upper triangular system of equations  $U.x = y$
- Finally solve for  $x$  using back substitution
- When this needs to be solved multiple times, usually perform LU -factorization using similar ideas.
- Serial runtime takes  $n^2/2$  divisions and  $(n^3/3) - (n^2/2)$  subtractions and multiplications, total approximately,  $W = \frac{2}{3}n^3$

# Serial Gaussian Elimination Algorithm

---

---

```
1.  procedure GAUSSIAN_ELIMINATION ( $A, b, y$ )
2.  begin
3.    for  $k := 0$  to  $n - 1$  do           /* Outer loop */
4.    begin
5.      for  $j := k + 1$  to  $n - 1$  do
6.         $A[k, j] := A[k, j] / A[k, k];$  /* Division step */
7.       $y[k] := b[k] / A[k, k];$ 
8.       $A[k, k] := 1;$ 
9.      for  $i := k + 1$  to  $n - 1$  do
10.     begin
11.       for  $j := k + 1$  to  $n - 1$  do
12.          $A[i, j] := A[i, j] - A[i, k] \times A[k, j];$  /* Elimination step */
13.        $b[i] := b[i] - A[i, k] \times y[k];$ 
14.        $A[i, k] := 0;$ 
15.     endfor;           /* Line 9 */
16.   endfor;           /* Line 3 */
17. end GAUSSIAN_ELIMINATION
```

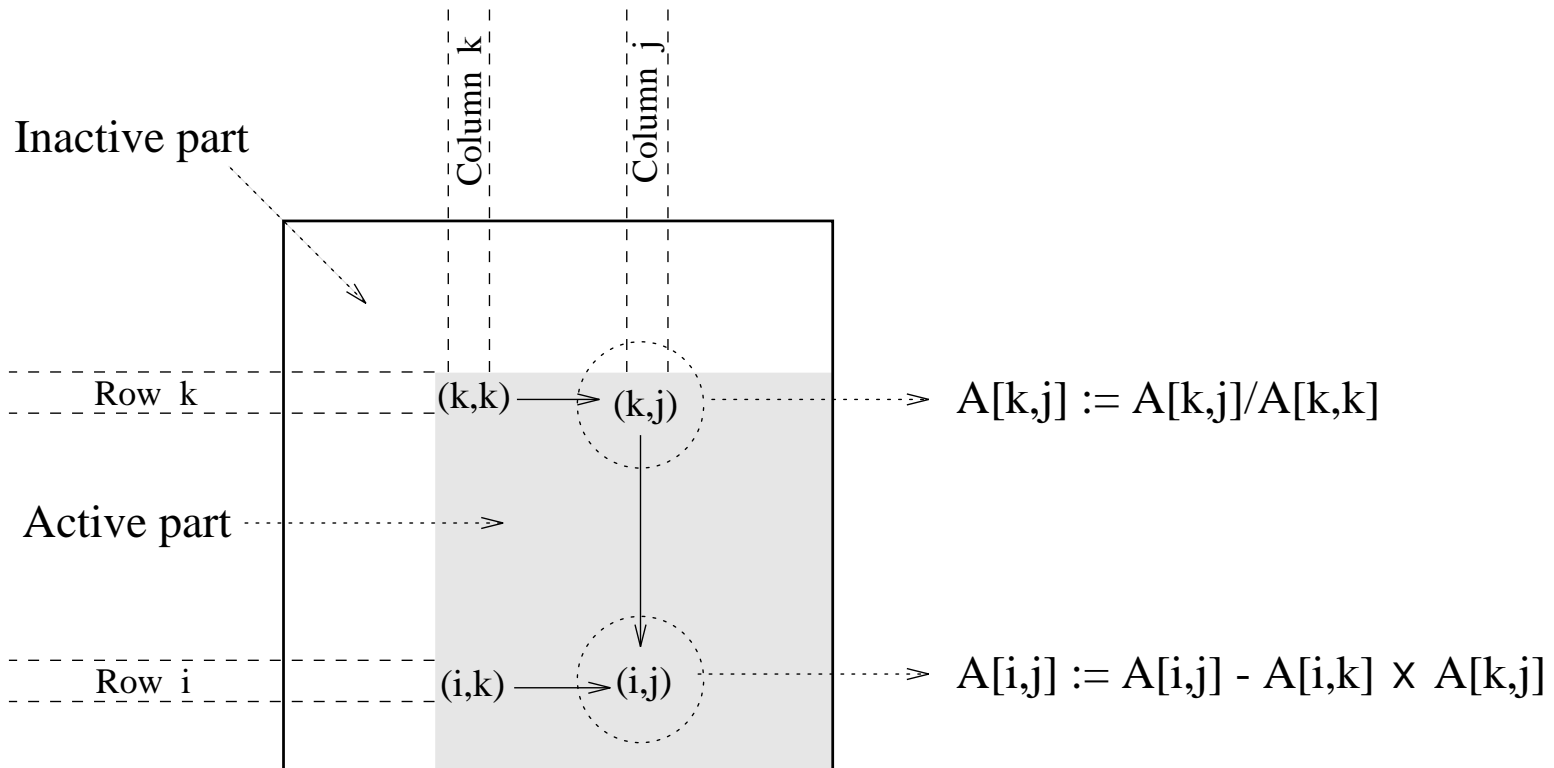
---

**Program 5.4** A serial Gaussian elimination algorithm that converts the system of linear equations  $Ax = b$  to a unit upper-triangular system  $Ux = y$ . The matrix  $U$  occupies the upper-triangular locations of  $A$ . This algorithm assumes that  $A[k, k] \neq 0$  when it is used as a divisor on lines 6 and 7.

Copyright (r) 1994 Benjamin/Cummings Publishing Co.

# Serial Gaussian Elimination Algorithm-Figure

---



# Parallel Gaussian Elimination Algorithm

---

- Can use either two-dimensional blocked or cyclic distribution of matrix

1	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)	(0,7)
0	1	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)	(1,7)
0	0	1	(2,3)	(2,4)	(2,5)	(2,6)	(2,7)
0	0	0	(3,3)	(3,4)	(3,5)	(3,6)	(3,7)
0	0	0	(4,3)	(4,4)	(4,5)	(4,6)	(4,7)
0	0	0	(5,3)	(5,4)	(5,5)	(5,6)	(5,7)
0	0	0	(6,3)	(6,4)	(6,5)	(6,6)	(6,7)
0	0	0	(7,3)	(7,4)	(7,5)	(7,6)	(7,7)

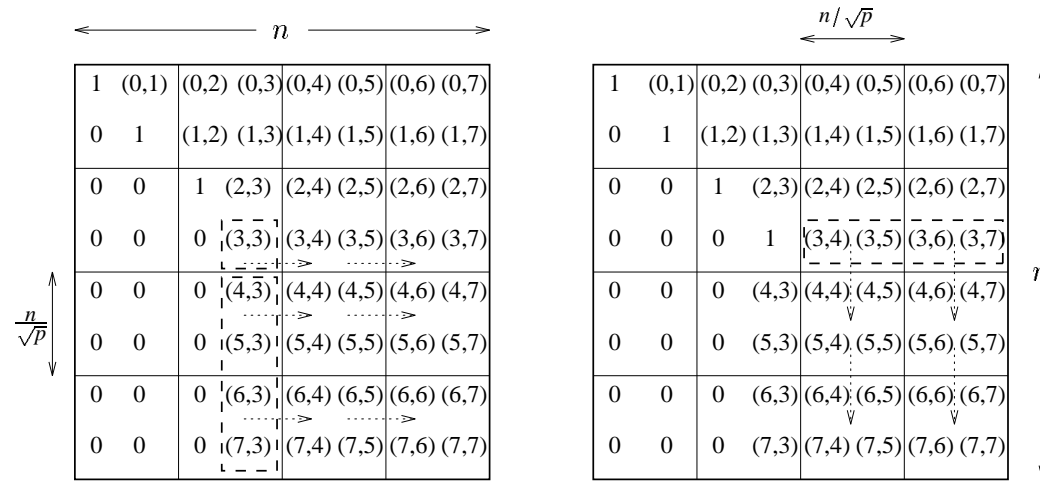
(a) Block-checkerboard mapping

1	(0,4)	(0,1)	(0,5)	(0,2)	(0,6)	(0,3)	(0,7)
0	(4,4)	0	(4,5)	0	(4,6)	(4,3)	(4,7)
0	(1,4)	1	(1,5)	(1,2)	(1,6)	(1,3)	(1,7)
0	(5,4)	0	(5,5)	0	(5,6)	(5,3)	(5,7)
0	(2,4)	0	(2,5)	1	(2,6)	(2,3)	(2,7)
0	(6,4)	0	(6,5)	0	(6,6)	(6,3)	(6,7)
0	(3,4)	0	(3,5)	0	(3,6)	(3,3)	(3,7)
0	(7,4)	0	(7,5)	0	(7,6)	(7,3)	(7,7)

(b) Cyclic-checkerboard mapping

# Parallel Gaussian Elimination Algorithm

- Communication steps in step  $k=3$  for  $8 \times 8$  matrix on 16 processors.



(a) Rowwise broadcast of  $A[i,k]$   
for  $i = k$  to  $(n - 1)$

(b) Columnwise broadcast of  $A[k,j]$   
for  $j = (k + 1)$  to  $(n - 1)$



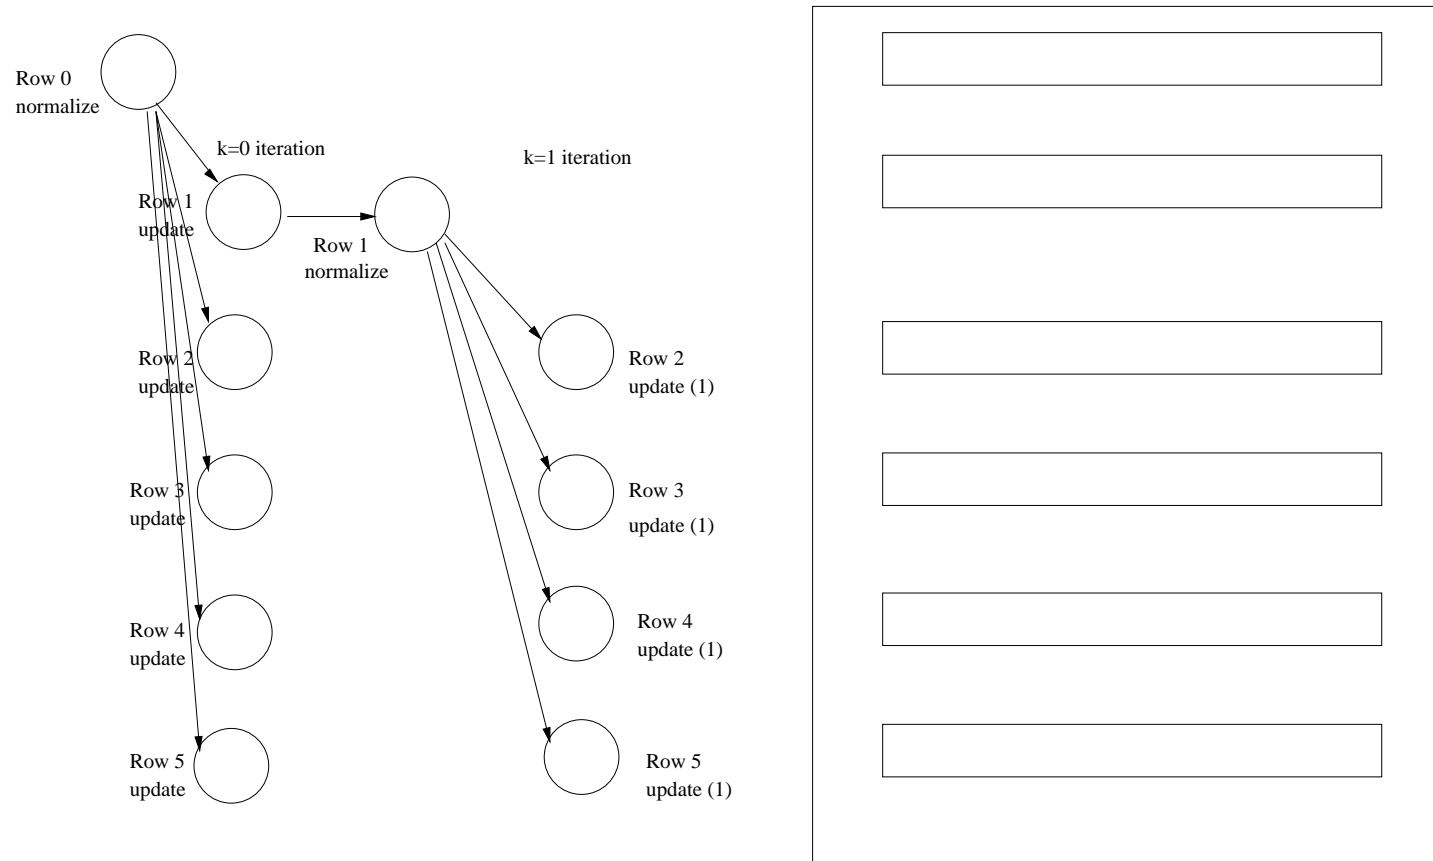
# Optimizations on Parallel Algorithm

---

- In previous algorithm, processors worked synchronously in each step.
- At any given time, all processors worked on the same iteration, i.e. the  $(k+1)$ st iteration started only after all the computation and communication of the  $k$ th iteration was complete.
- Can design an asynchronous, or pipelined algorithm, where some processors perform the computation of a given iteration earlier than other processors, and start working on subsequent iterations sooner.
- Can overlap computations with communication.

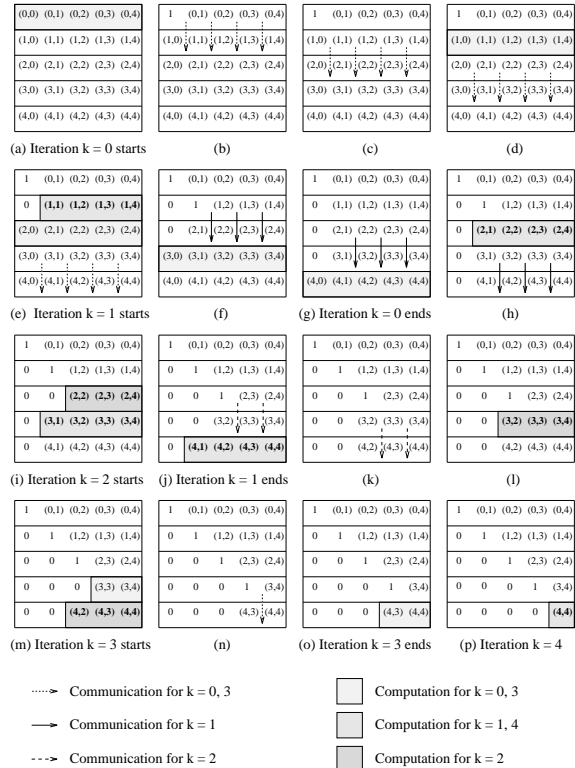
# Overview of Pipelined Parallel GE

---



# Illustration of Pipelined Parallel GE

- Assume that processor can do only one of compute/send/receive at a time



# Summary

---

- Design of parallel algorithms
- Dense matrix representations
- Matrix transpose algorithms
- Matrix vector multiplication
- Matrix matrix multiplication
- Linear system of equation solvers
- NEXT LECTURE: Sparse Matrix Algorithms